

Gradients manipulations

Optimization for Machine Learning — Exercise 01

Monday 17th April, 2023

Recall that the gradient of a differentiable function $f: \mathbb{R}^m \rightarrow \mathbb{R}$ at $x \in \mathbb{R}^m$ is a vector in \mathbb{R}^m , usually denoted by $\nabla f(x)$, such that

$$\nabla f(x) = \begin{pmatrix} \partial_{x_1} f(x) \\ \vdots \\ \partial_{x_m} f(x) \end{pmatrix}.$$

If f is multivalued, i.e. $f: \mathbb{R}^m \rightarrow \mathbb{R}^n$, then its *Jacobian* at x , denoted by $J_f(x)$, is the $n \times m$ matrix such that

$$J_f(x) = \left(\frac{\partial f(x)_i}{\partial x_j} \right)_{i=1, j=1}^{n, m} = \begin{pmatrix} \frac{\partial f(x)_1}{\partial x_1} & \frac{\partial f(x)_1}{\partial x_2} & \dots & \frac{\partial f(x)_1}{\partial x_m} \\ \frac{\partial f(x)_2}{\partial x_1} & \ddots & & \vdots \\ \vdots & & & \frac{\partial f(x)_n}{\partial x_m} \end{pmatrix}.$$

When $n = 1$, we therefore have $\nabla f(x) = (J_f(x))^\top$.

Product rule Similar to the 1D case, a product rule exists for multivariate functions.

1. a) If $f: \mathbb{R}^d \rightarrow \mathbb{R}^k$ and $g: \mathbb{R}^d \rightarrow \mathbb{R}^k$, show that $\nabla(f^\top g)(x) = J_f(x)^\top g(x) + J_g(x)^\top f(x)$.
b) What happens with $k = 1$?

Chain rule The chain rule for Jacobian is, for $f: \mathbb{R}^m \rightarrow \mathbb{R}^k$, and $g: \mathbb{R}^k \rightarrow \mathbb{R}^n$, $J_{g \circ f}(x) = J_g(f(x))J_f(x)$ (when the composition makes sense, and everything is differentiable). Compare with the chain rule in the 1D case: $(g \circ f)'(x) = g'(f(x)) \cdot f'(x)$.

2. a) Compute the gradient of $g \circ f$ at $x \in \mathbb{R}^m$, when $f: \mathbb{R}^m \rightarrow \mathbb{R}^k$ and $g: \mathbb{R}^k \rightarrow \mathbb{R}$, as a function of $J_f(x)$ and $\nabla g(f(x))$ (note the difference between $\nabla(g \circ f)(x)$ and $\nabla g(f(x))$).
b) Compute the gradient of $h_2 \circ h_1 \circ f$, when $h_1: \mathbb{R} \rightarrow \mathbb{R}$ and $h_2: \mathbb{R} \rightarrow \mathbb{R}$ are single valued functions, and $f: \mathbb{R}^d \rightarrow \mathbb{R}$ is a vector function.

- c) Assume that $x(w), y(w), z(w)$ are function of $w \in \mathbb{R}^p$, and that $\mathcal{L}: \mathbb{R}^3 \rightarrow \mathbb{R}$ is a function of x, y, z . Show that $\nabla_w \mathcal{L}(x(w), y(w), z(w)) = \frac{\partial \mathcal{L}(x, y, z)}{\partial x} \nabla x(w) + \frac{\partial \mathcal{L}(x, y, z)}{\partial y} \nabla y(w) + \frac{\partial \mathcal{L}(x, y, z)}{\partial z} \nabla z(w)$.

Classical vector functions Often, the functions that will appear are build from simpler ones, such as the linear product $\langle a, b \rangle = a^\top b$, or a matrix-vector multiplication $A \cdot b$, etc. The easiest to find the gradient of such function is usually to go back to the expression with the indices, e.g. $\langle a, b \rangle = \sum_i a_i b_i$. Another method is to look up if the formula is in the Matrix Cookbook.

3. a) Let $a \in \mathbb{R}^n$. Show that $\nabla_x \langle a, x \rangle = a$,
 b) Let $A \in \mathbb{R}^{n \times m}$, and $f: \mathbb{R}^m \rightarrow \mathbb{R}^n, x \mapsto Ax$. Show that $J_f(x) = A$.
 c) What is $\nabla_x \langle x, Ax \rangle$?
 d) What is $\nabla_A \langle x, Ax \rangle$?

General functions

4. Compute the gradients of the functions:
 a) $f: \mathbb{R}^3 \rightarrow \mathbb{R}, (x, y, z) \mapsto \frac{1}{2}x^2 + yz - \ln(1 + \exp(x^2 y^3 z))$
 b) $g: \mathbb{R}^d \rightarrow \mathbb{R}, x \mapsto \frac{1}{2}\|x\|^2 = \frac{1}{2}x^\top x = \frac{1}{2} \sum_{i=1}^d x_i^2$
 c) $h: \mathbb{R}^d \times \mathbb{R}^d \rightarrow \mathbb{R}, (x, w) \mapsto \ln(1 + \exp(-w^\top x))$. Compute the gradient with respect to w : $\nabla_w h(x, w)$.